

Lab - PostGIS



PostgreSQL

- ORDBMS
- Open Source
- Easy to add custom functions
- Support many third-party library



PostgreSQL

What is PostGIS



PostgreSQL

+



GIS

=



PostGIS

(spatial types, indexes, and functions)

Other Spatial DB

- Spatialite
- Oracle spatial
- MySQL/MariaDB spatial
- Microsoft SQLServer spatial database



Installation Guide

- For Windows -
(Mac OS is later introduced)

PostgreSQL Installation

- Find the corresponding OS version From <https://www.enterprisedb.com/downloads/postgres-postgresql-downloads>
- ☆ Install version 11.x

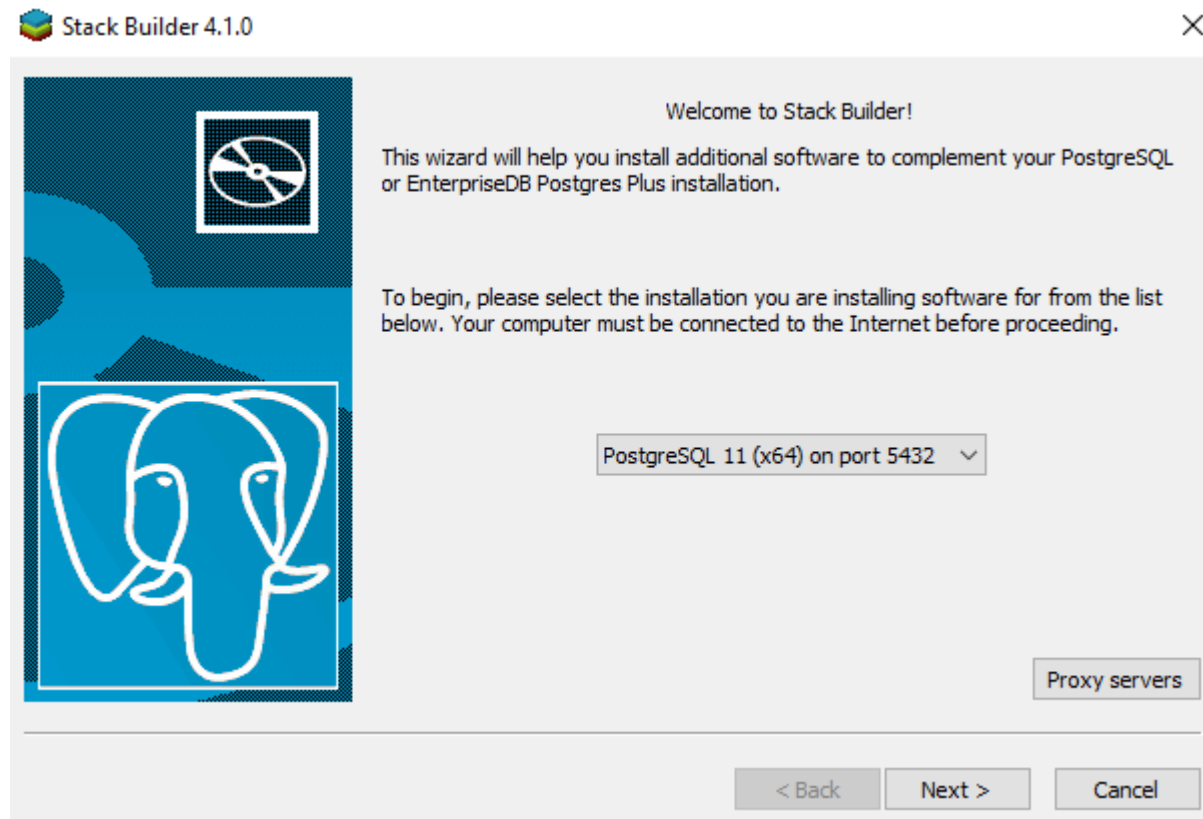
PostgreSQL Version	Linux x86-64	Linux x86-32	Mac OS X	Windows x86-64	Windows x86-32
11.0	N/A	N/A	Download	Download	N/A
10.5	Download	Download	Download	Download	Download
9.6.10	Download	Download	Download	Download	Download
9.5.14	Download	Download	Download	Download	Download
9.4.19	Download	Download	Download	Download	Download
9.3.24	Download	Download	Download	Download	Download

PostgreSQL Installation

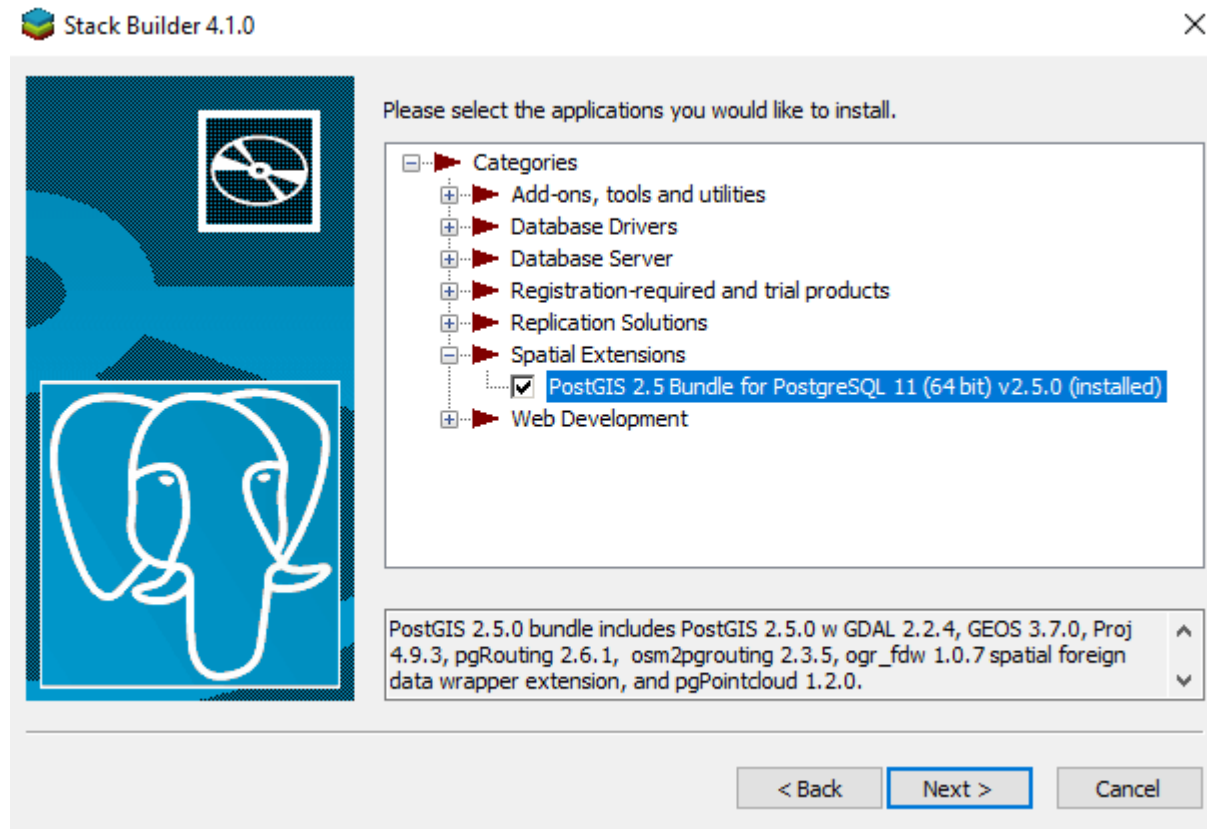
- Follow the installation guide

https://www.enterprisedb.com/docs/en/11.0/PG_Inst_Guide_v11/PostgreSQL_Installation_Guide.1.08.html

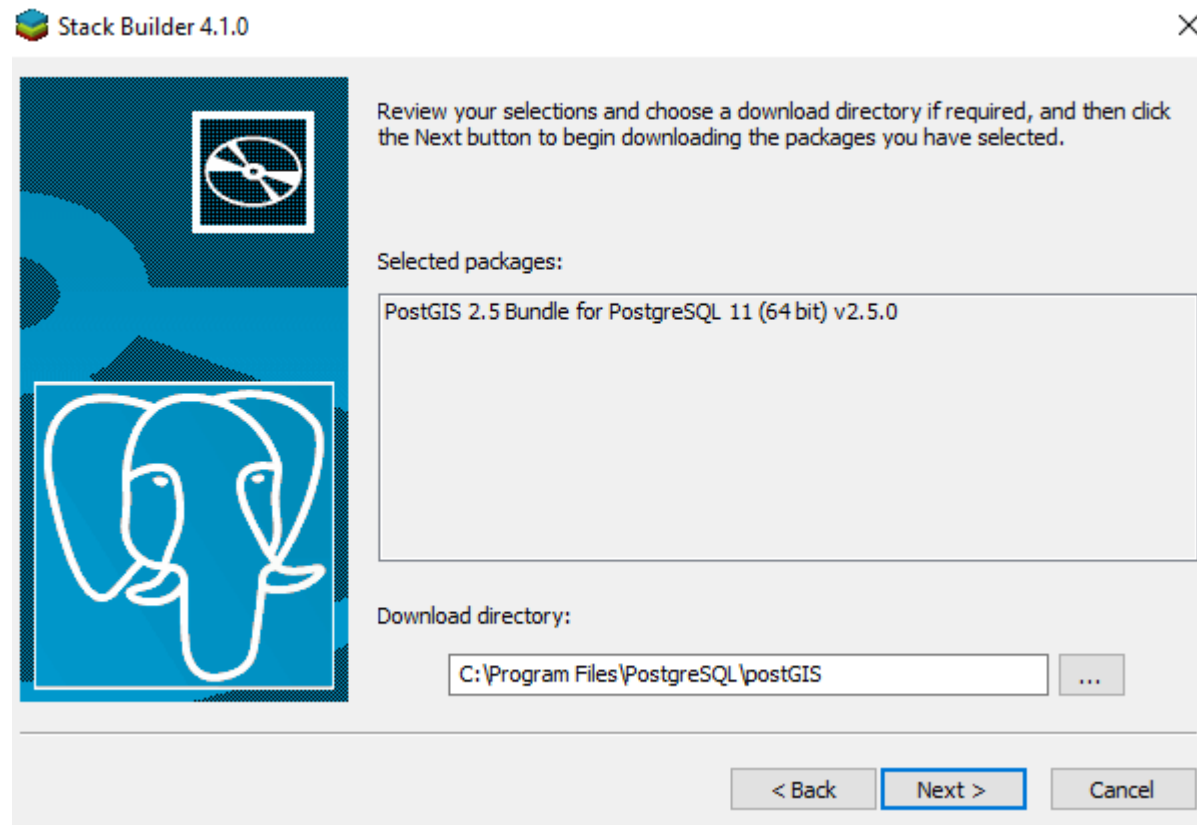
Stack Builder



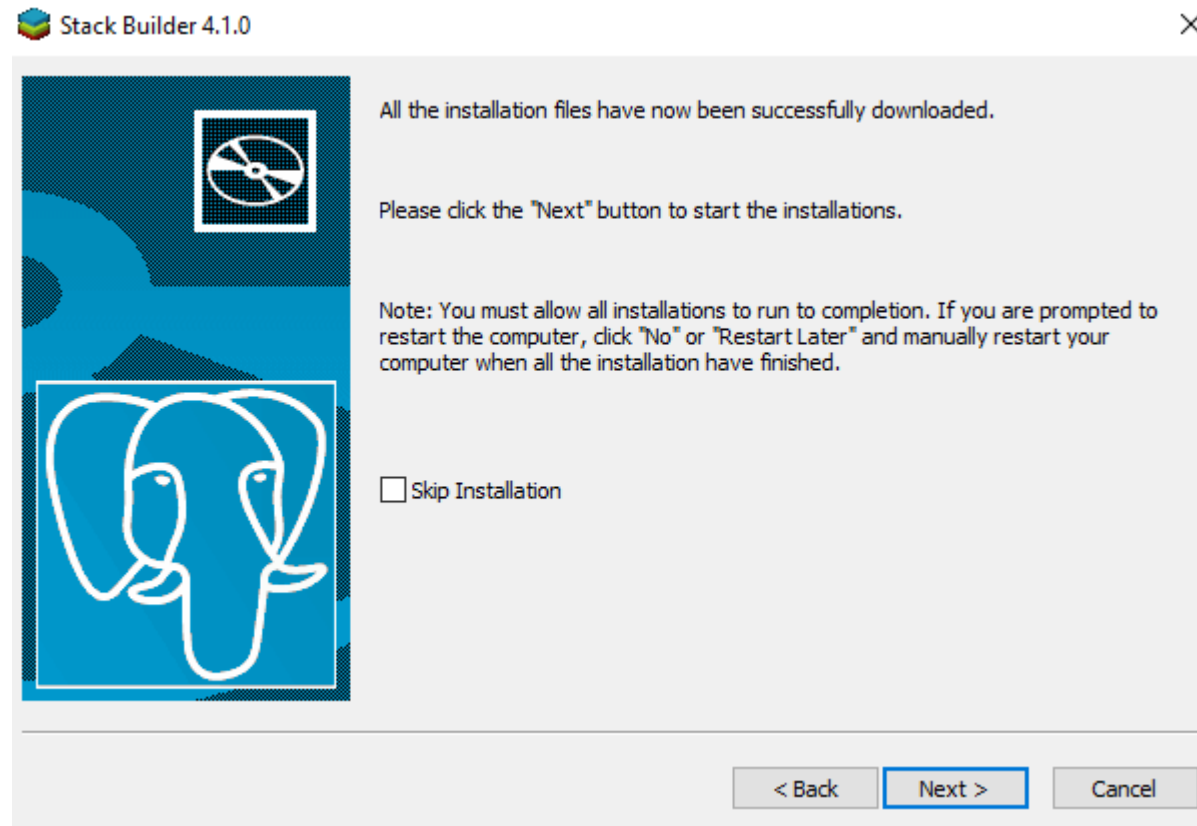
Stack Builder



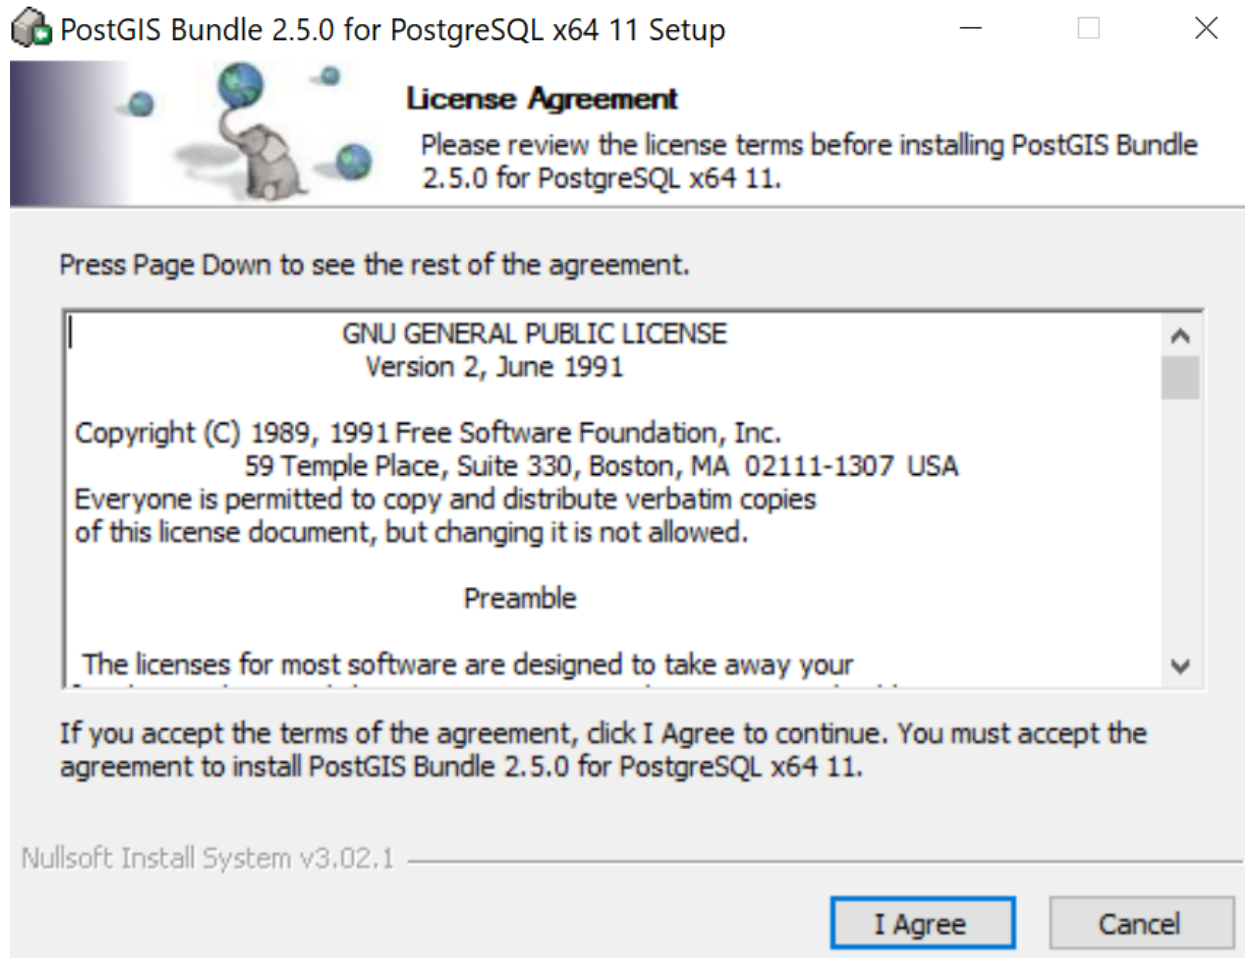
Stack Builder



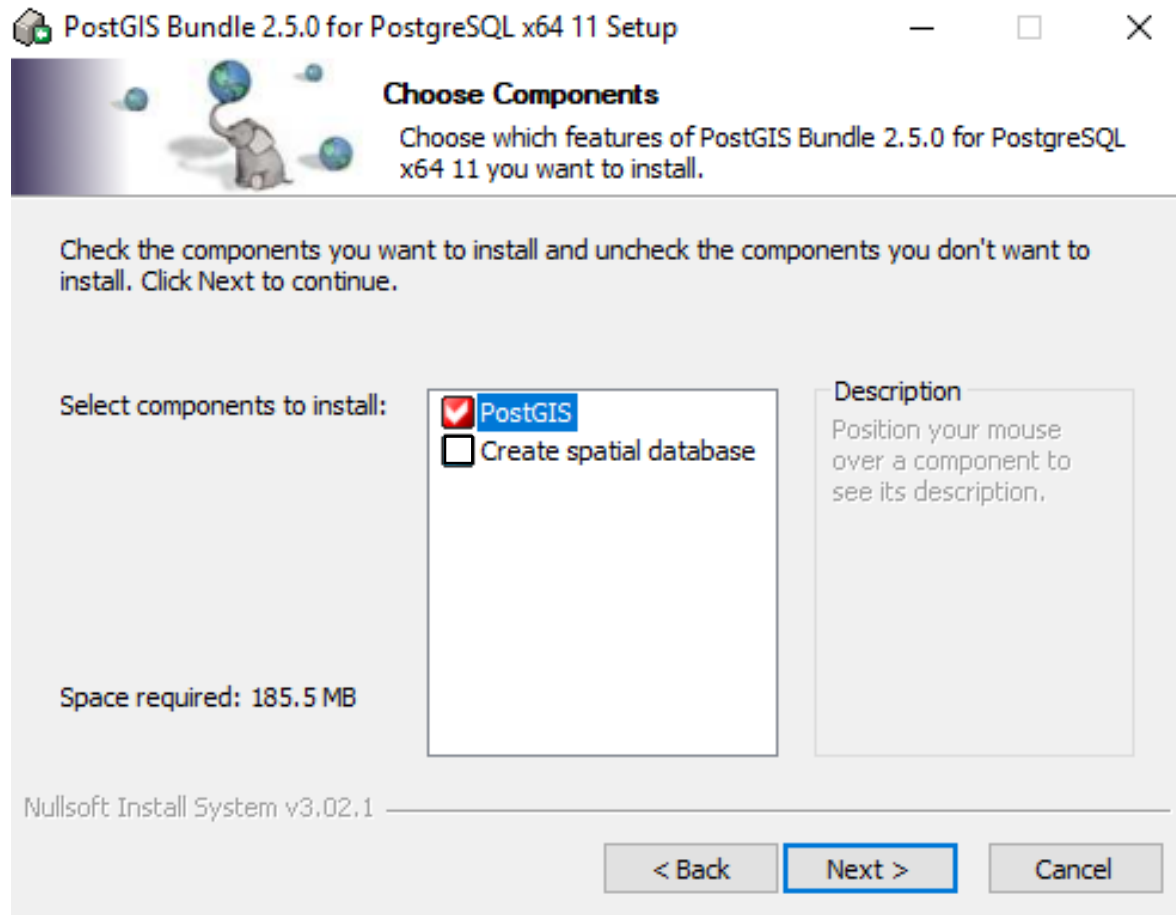
Stack Builder



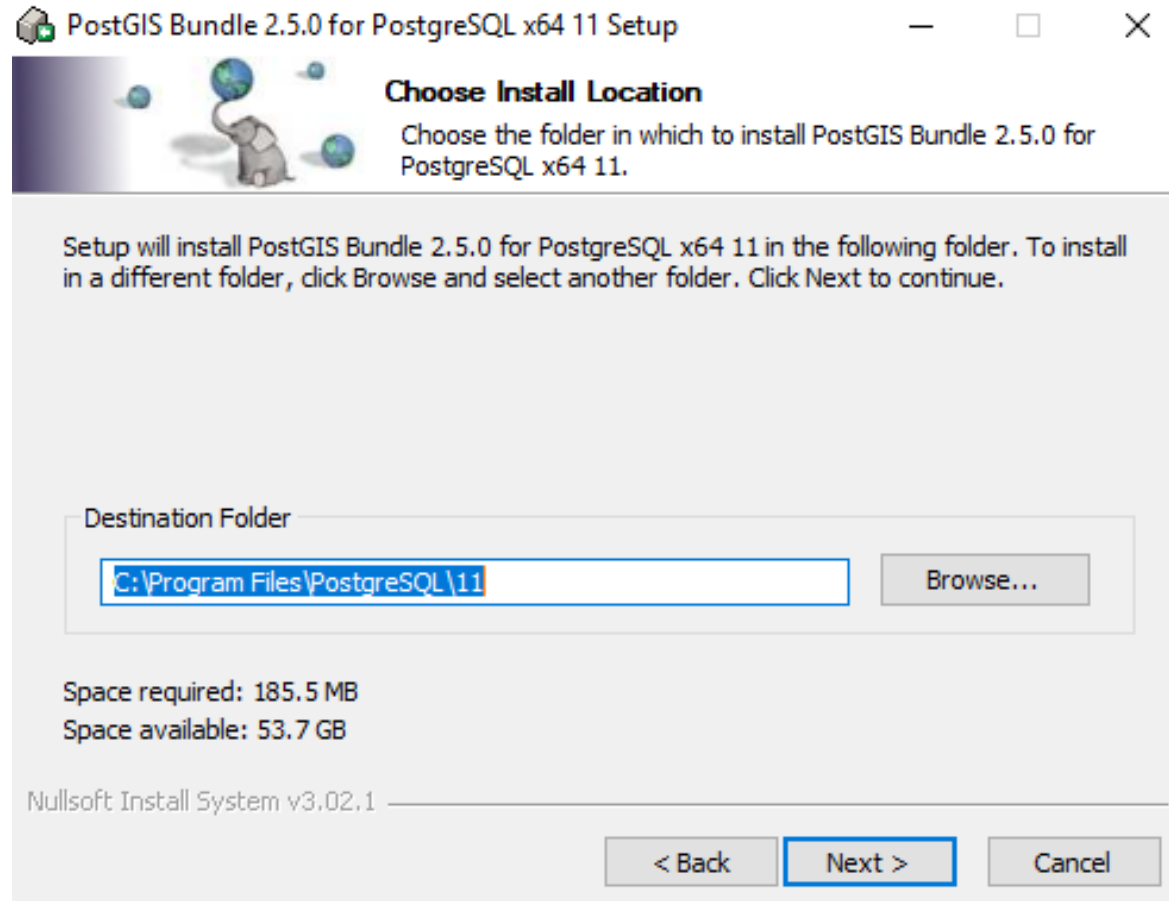
PostGIS Installer



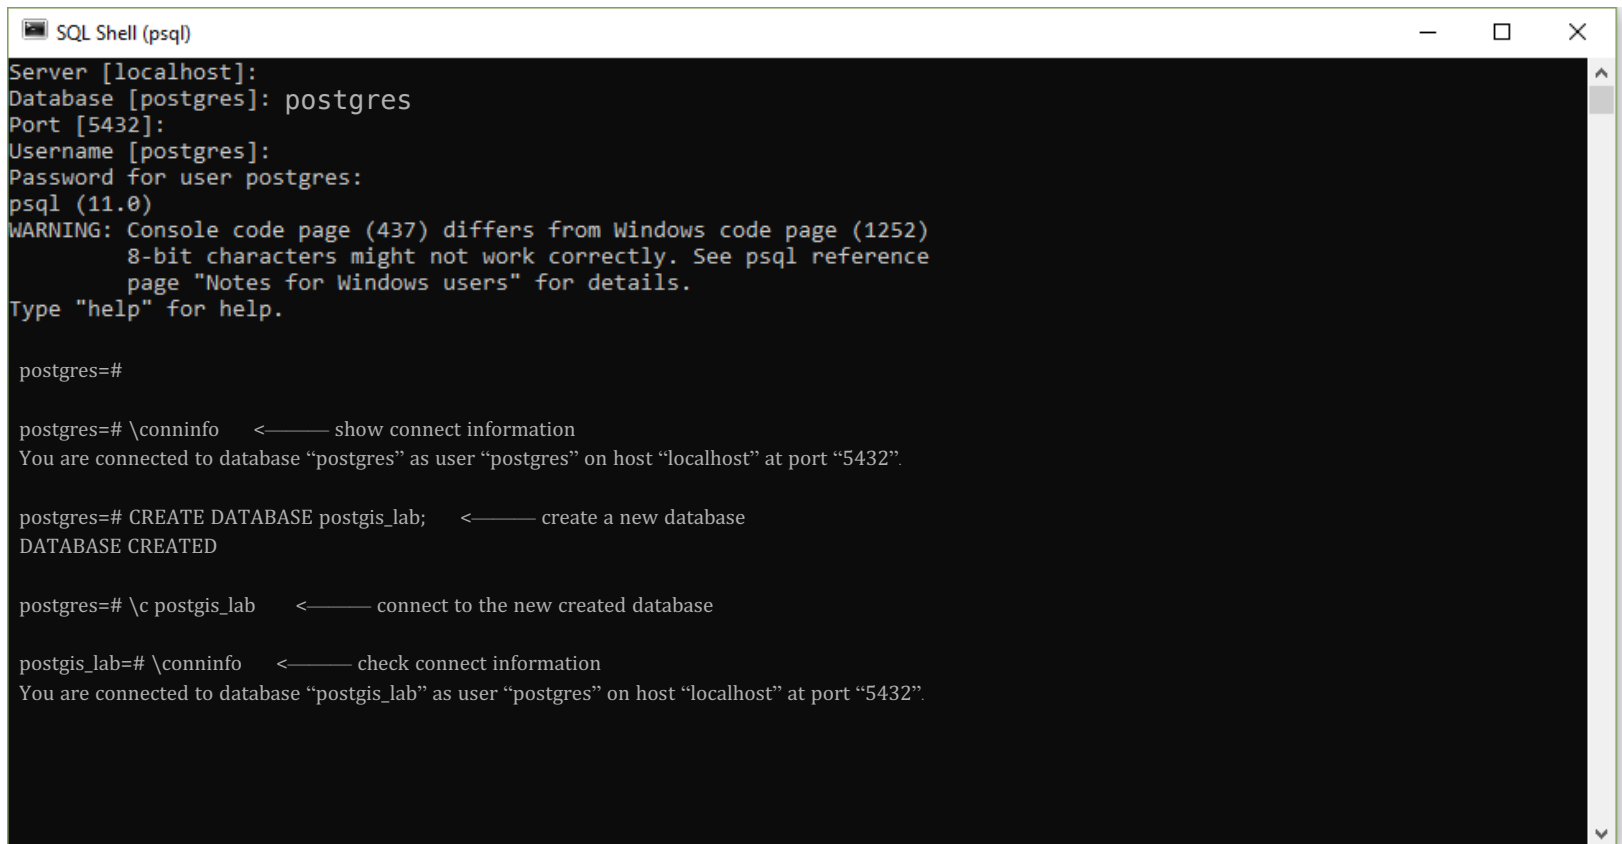
PostGIS Installer



PostGIS Installer



After installing, open psql



```
SQL Shell (psql)
Server [localhost]:
Database [postgres]: postgres
Port [5432]:
Username [postgres]:
Password for user postgres:
psql (11.0)
WARNING: Console code page (437) differs from Windows code page (1252)
         8-bit characters might not work correctly. See psql reference
         page "Notes for Windows users" for details.
Type "help" for help.

postgres=#

postgres=# \conninfo    <----- show connect information
You are connected to database "postgres" as user "postgres" on host "localhost" at port "5432".

postgres=# CREATE DATABASE postgis_lab;    <----- create a new database
DATABASE CREATED

postgres=# \c postgis_lab    <----- connect to the new created database

postgis_lab=# \conninfo    <----- check connect information
You are connected to database "postgis_lab" as user "postgres" on host "localhost" at port "5432".
```

Installation Guide

- For Mac OS -

Homebrew Installation

- Website: <https://brew.sh/>

Install Homebrew

```
/usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```

Paste that in a macOS Terminal prompt.

The script explains what it will do and then pauses before it does it. Read about other **installation options**. Install Homebrew on **Linux and Windows Subsystem for Linux**.

Postgres & PostGIS Installation

- Remove old versions (Optional)

```
$ brew uninstall --force postgis postgresql  
$ rm -rf /usr/local/var/postgres
```

- Install new ones

```
$ brew install postgres postgis
```

- Start server

```
$ pg_ctl -D /usr/local/var/postgres start
```

- Shutdown server

```
$ pg_ctl -D /usr/local/var/postgres -l logfile stop
```

Post installation

- Create database storage

```
$ initdb /usr/local/var/postgres
```

- If terminal shows an error

```
initdb: directory "/usr/local/var/postgres" exists but is not empty
If you want to create a new database system, either remove or empty
the directory "/usr/local/var/postgres" or run initdb
with an argument other than "/usr/local/var/postgres".
```

- Remove old database files `$ rm -r /usr/local/var/postgres`
- Re-run the initdb command `$ initdb /usr/local/var/postgres`

- Create a new database

```
$ createdb postgis_lab
```

- Operate the created database with psql

```
$ psql postgis_lab
```

```
→ postgis_lab=#
```

PSQL

SQL shell for PostgreSQL

PSQL

- Useful commands in psql (`postgres_lab=#`)
 - Check for all command `\?`
 - List table, view `\d`
 - Connect to another database `\c {DBNAME}`
 - Change the current working directory `\cd [DIR]`
 - Quit `\q` or CTRL+D

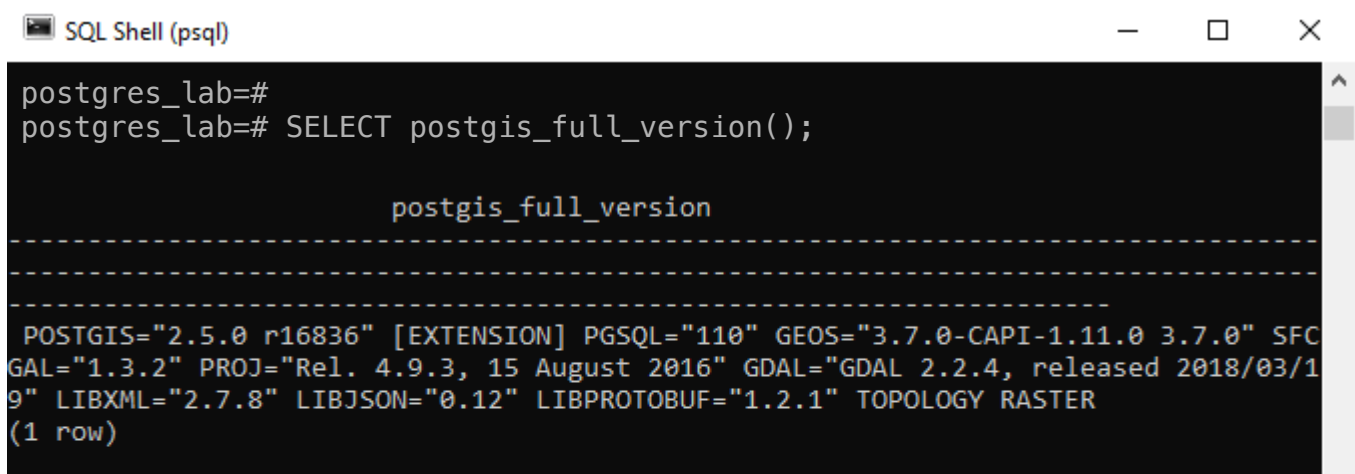
Enable GIS Function in Postgres

- Load PostGIS spatial extension

```
# CREATE EXTENSION postgis;
```

- Confirm whether PostGIS installed successfully

```
# SELECT postgis_full_version();
```



```
SQL Shell (psql)
postgres_lab=#
postgres_lab=# SELECT postgis_full_version();

               postgis_full_version
-----
POSTGIS="2.5.0 r16836" [EXTENSION] PGSQL="110" GEOS="3.7.0-CAPI-1.11.0 3.7.0" SFC
GAL="1.3.2" PROJ="Rel. 4.9.3, 15 August 2016" GDAL="GDAL 2.2.4, released 2018/03/1
9" LIBXML="2.7.8" LIBJSON="0.12" LIBPROTOBUF="1.2.1" TOPOLOGY RASTER
(1 row)
```

Loading Spatial Data

Spatial Data format

- **Shapefile:** a popular vector data for GIS format
- **GeoJSON:** JSON format for the spatial extended specification
- **GML:** XML format for representing spatial feature information.
- **KML:** the spatial XML format used by Google Earth

Shapefiles

- **.shp** : shape format; the feature geometry itself
- **.shx** : shape index format; a positional index of the feature geometry
- **.dbf** : attribute format; columnar attributes for each shape
- **(.prj)** : projection format; the coordinate system and projection information

SRID(Spatial Reference Identifier)

- It packs all the information about a map projection (.prj) into a single number.

```
PROJCS["NAD83 / UTM zone 18N",  
  GEOGCS["NAD83",  
    DATUM["North_American_Datum_1983",  
      SPHEROID["GRS 1980",6378137,298.257222101,AUTHORITY["EPSG","7019"]],  
      AUTHORITY["EPSG","6269"]],  
    PRIMEM["Greenwich",0,AUTHORITY["EPSG","8901"]],  
    UNIT["degree",0.01745329251994328,AUTHORITY["EPSG","9122"]],  
    AUTHORITY["EPSG","4269"]],  
  UNIT["metre",1,AUTHORITY["EPSG","9001"]],  
  PROJECTION["Transverse_Mercator"],  
  PARAMETER["latitude_of_origin",0],  
  PARAMETER["central_meridian",-75],  
  PARAMETER["scale_factor",0.9996],  
  PARAMETER["false_easting",500000],  
  PARAMETER["false_northing",0],  
  AUTHORITY["EPSG","26918"],  
  AXIS["Easting",EAST],  
  AXIS["Northing",NORTH]]
```

SRID (Spatial Reference Identifier)

- Plug the contents of the .prj file into <http://prj2epsg.org>.
- This will give you the number that most closely match your projection definition.

```
$ python3 esriprj2standards.py [prj file]
```

Loading data into database (1)

- Go to cmd/terminal
- Use shp2pgsql tool to convert shapefile into database.

```
$ shp2pgsql -s SRID FILE.shp | psql -d DBNAME -U USERNAME
```

- (Find Username)

```
> SELECT username from pg_user;
```

```
postgres=# SELECT username from pg_user;
username
-----
yenhao
(1 row)
```

```
$ shp2pgsql -s 26918 nyc_streets.shp | psql -d postgres -U postgres
```

```
-----
INSERT INTO "nyc_streets" ("id","name","oneway","type",geom) VALUES ('1020','FDR
Dr',NULL,'residential','01050000202669000001000000010200000002000000C9D2C4E351F52
141E91561A27037514170DAEFAB3FF52141C419EE3D72375141');
INSERT INTO "nyc_streets" ("id","name","oneway","type",geom) VALUES ('1152','FDR
Dr',NULL,'residential','010500002026690000010000000102000000020000004F6C8F179AF72
141E23CDE3EEB375141B95D381AFAF721417829E25FFE375141');
```

For Windows

Windows PowerShell

```
PS C:\Users\refu\Desktop\data> shp2pgsql -s 26918 .\nyc_streets.shp | psql -d postgres -U postgres
Field id is an FTDoube with width 11 and precision 0
Shapefile type: Arc
Postgis type: MULTILINESTRING[2]
Password for user postgres:
```

Loading data into database (2)

- After data loaded

```
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
COMMIT
ANALYZE
```

- Check data inserted by


> \d

```
postgres_lab=# \d
               List of relations
Schema |      Name      | Type   | Owner
-----+-----+-----+-----
public | geography_columns | view   | yenhao
public | geometry_columns | view   | yenhao
public | nyc_streets      | table  | yenhao
public | nyc_streets_gid_seq | sequence | yenhao
public | spatial_ref_sys  | table  | yenhao
(5 rows)
```

- Repeat for the rest of data
(nyc_subway_stations/ nyc_neighborhoods/
nyc_homicides/ nyc_census_blocks.shx)

nyc_census_blocks

- A census block is the smallest geography for which census data is reported



blkid

popn_total

popn_white

popn_black

popn_nativ

popn_asian


popn_other

boroname

geom

nyc_neighborhoods

- Neighborhoods are social constructs that do not follow lines laid down by the government.




name

boroname

geom

nyc_streets



name

oneway

type

geom

nyc_subway_stations



name

borough

routes

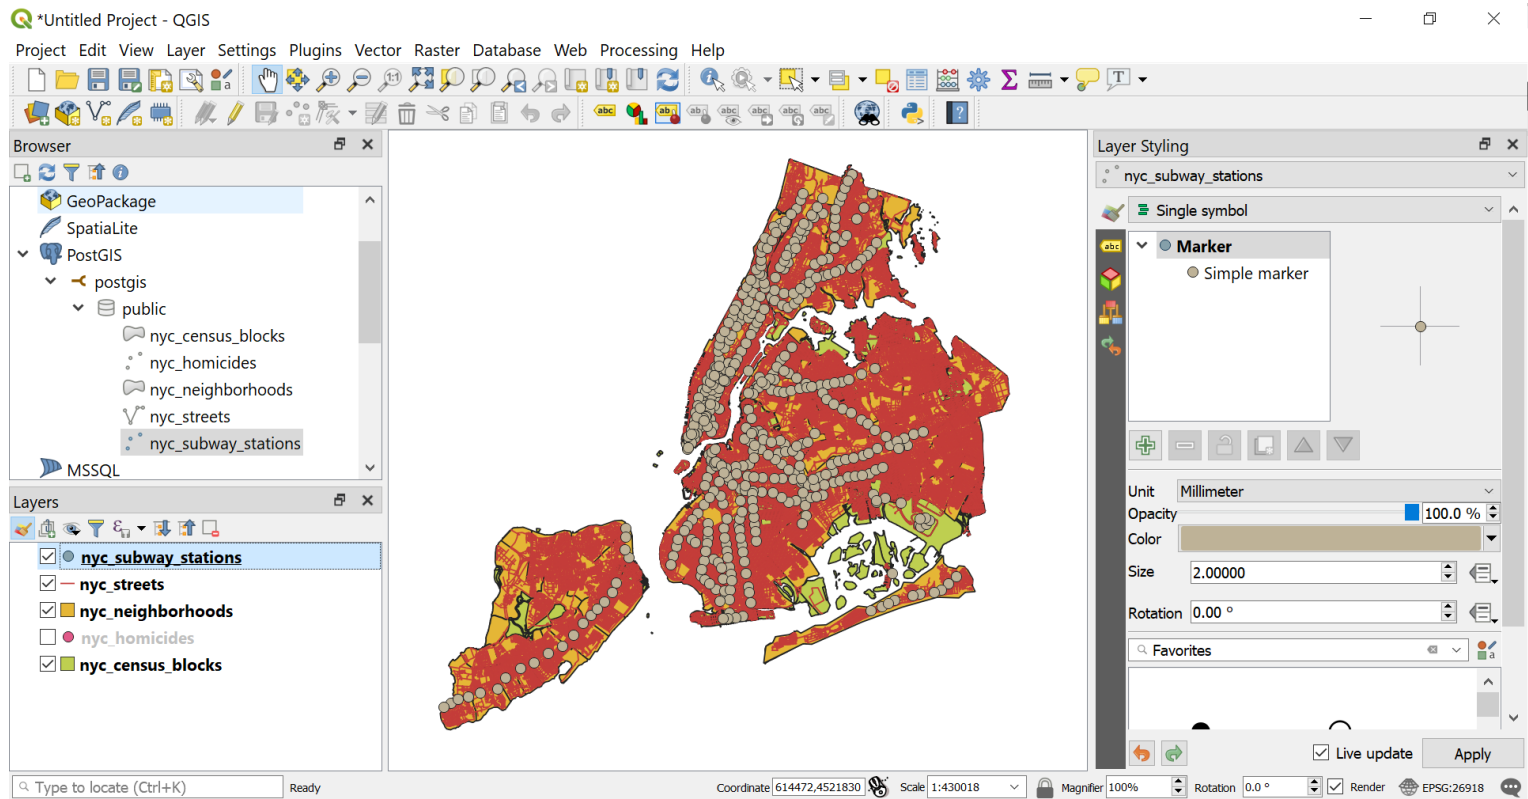
transfers

express

geom

QGIS

- You can use QGIS to visualize data.



Spatial SQL

Basic PostgreSQL

- Find column_names of table

```
postgres=# SELECT
postgres=#     COLUMN_NAME
postgres=# FROM
postgres=#     information_schema.COLUMNS
postgres=# WHERE
postgres=#     TABLE_NAME = 'nyc_neighborhoods';
 column_name
-----
gid
boroname
name
geom
(4 rows)
```

- Or
> \d table_name

Basic PostgreSQL

```
> SELECT name
   FROM nyc_neighborhoods
  WHERE boroname = 'Brooklyn';
```

```
postgis=# SELECT name
postgis-#       FROM nyc_neighborhoods
postgis-#       WHERE boroname = 'Brooklyn';
       name
-----
Bensonhurst
Bay Ridge
Boerum Hill
Cobble Hill
Downtown
Sunset Park
Borough Park
East Brooklyn
Flatbush
Park Slope
Williamsburg
Canarsie
Greenwood
Gravesend-Sheepshead Bay
Dyker Heights
Brownsville
Bushwick
Fort Green
Mapleton-Flatlands
Bedford-Stuyvesant
Carroll Gardens
Coney Island
Red Hook
(23 rows)
```

Basic PostgreSQL - exercise

- For each borough, what percentage of the population is white?
- Table : **nyc_census_blocks**

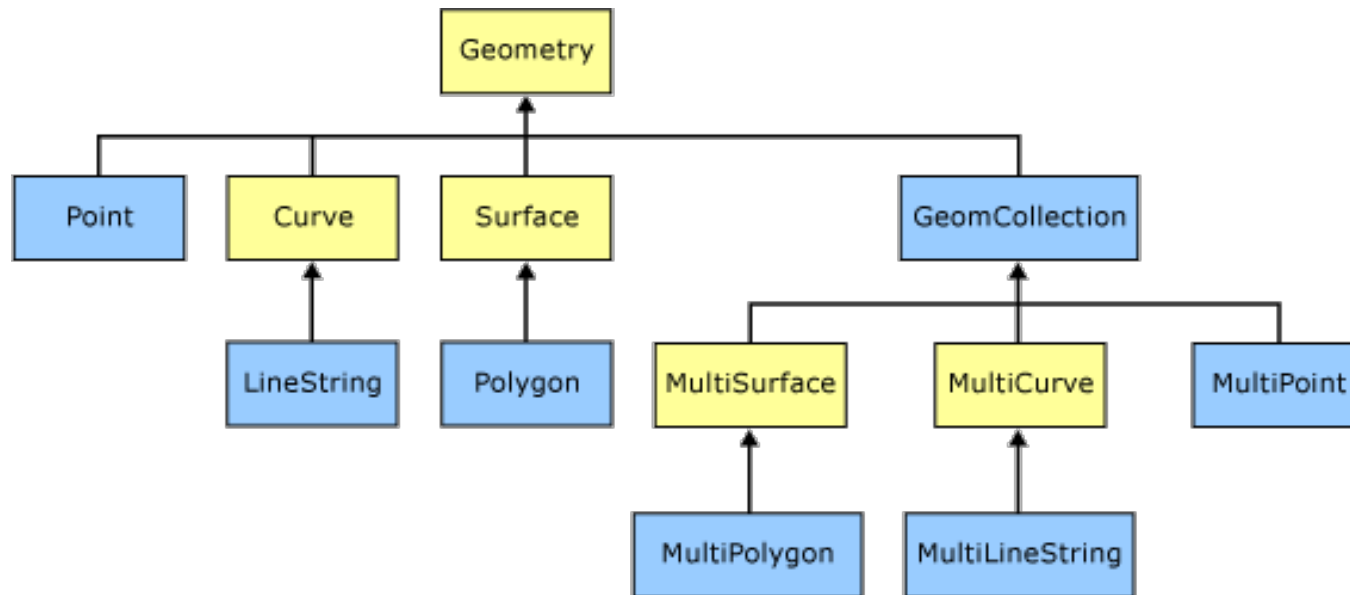
A census block is the smallest geography for which census data is reported. All higher level census geographies (block groups, tracts, metro areas, counties, etc) can be built from unions of census blocks. We have attached some demographic data to our collection of blocks.

Number of records: 36592

blkid	A 15-digit code that uniquely identifies every census block . Eg: 360050001009000
popn_total	Total number of people in the census block
popn_white	Number of people self-identifying as “White” in the block
popn_black	Number of people self-identifying as “Black” in the block
popn_nativ	Number of people self-identifying as “Native American” in the block
popn_asian	Number of people self-identifying as “Asian” in the block
popn_other	Number of people self-identifying with other categories in the block
boroname	Name of the New York borough. Manhattan, The Bronx, Brooklyn, Staten Island, Queens
geom	Polygon boundary of the block

Spatial Data Types - Geometry

- Point, LineString, Polygon, GeoCollection



Geometry - create

```
> CREATE TABLE geometries (name varchar, geom geometry);

> INSERT INTO geometries VALUES
  ('Point', 'POINT(0 0)'),
  ('Linestring', 'LINESTRING(0 0, 1 1, 2 1, 2 2)'),
  ('Polygon', 'POLYGON((0 0, 1 0, 1 1, 0 1, 0 0))'),
  ('PolygonWithHole', 'POLYGON((0 0, 10 0, 10 10, 0 10, 0 0),(1 1, 1 2,
2 2, 2 1, 1 1))'),
  ('Collection', 'GEOMETRYCOLLECTION(POINT(2 0),POLYGON((0 0, 1 0, 1 1,
0 1, 0 0)))');
```


Geometry - type

- Geometry type stores in **WKB(Well-known Binary)** format

```
postgis_lab=# SELECT * from geometries WHERE name='Point';
 name | geom
-----+-----
 Point | 0101000000000000000000000000000000000000000000000000000000000000
(1 row)
```

- **ST_AsText(geom)** can show in **WKT(Well-known Text)** format

```
postgis_lab=# SELECT name, ST_AsText(geom) from geometries WHERE name='Point';
 name | st_astext
-----+-----
 Point | POINT(0 0)
(1 row)

postgis_lab=# SELECT name, ST_AsText(geom) from geometries;
 name | st_astext
-----+-----
 Point | POINT(0 0)
 Linestring | LINESTRING(0 0,1 1,2 1,2 2)
 Polygon | POLYGON((0 0,1 0,1 1,0 1,0 0))
 PolygonWithHole | POLYGON((0 0,10 0,10 10,0 10,0 0),(1 1,1 2,2 2,2 1,1 1))
 Collection | GEOMETRYCOLLECTION(POINT(2 0),POLYGON((0 0,1 0,1 1,0 1,0 0)))
(5 rows)
```

Geometry - Point

- A spatial **point** represents a single location. Points are used to represent objects when details are not important at the target scale




- Useful spatial function
 - **ST_X(geom)** returns the X ordinate
 - **ST_Y(geom)** returns the Y ordinate
 - **ST_NDims(geom)** returns the number of dimensions
 - **Exercise: print out any station from “nyc_subway_stations”**

name	st_astext
Cortlandt St	POINT(583521.854408956 4507077.86259909)
(1 row)	

Geometry - Linestrings

- A **linestring** is a path between locations. It takes the form of an ordered series of two or more points.



- Useful spatial function
 - **ST_Length(geom)** returns the length of the linestring
 - **ST_StartPoint(geom)** returns the first coordinate as a point
 - **ST_EndPoint(geom)** returns the last coordinate as a point
 - **ST_NPoints(geom)** returns the number of coordinates in the linestring

```
postgis_lab=# SELECT name, ST_Length(geom),
postgis_lab=# ST_AsText(ST_StartPoint(geom)),
postgis_lab=# ST_AsText(ST_EndPoint(geom)), ST_NPoints(geom)
postgis_lab=# FROM geometries WHERE name = 'Linestring';
name | st_length | st_astext | st_astext | st_npoints
-----+-----+-----+-----+-----
Linestring | 3.414213562373095 | POINT(0 0) | POINT(2 2) | 4
(1 row)
```

Geometry - Polygons

- A **polygon** is a representation of an area. It is composed by rings.



- Useful spatial function
 - **ST_Area(geom)** returns the area of the polygons
 - **ST_NRings(geom)** returns the number of rings
 - **ST_ExteriorRing(geom)** returns the outer ring as a linestring
 - **ST_InteriorRingN(geom,n)** returns a specified interior ring as a linestring

```
postgis_lab=# SELECT name, ST_Area(geom), ST_NRings(geom),  
ST_AsText(ST_ExteriorRing(geom))  
FROM geometries WHERE name LIKE 'Polygon%';
```

name	st_area	st_nrings	st_astext
Polygon	1	1	LINESTRING(0 0,1 0,1 1,0 1,0 0)
PolygonWithHole	99	2	LINESTRING(0 0,10 0,10 10,0 10,0 0)

(2 rows)

- **Exercise: Please print the interior ring of “PolygonWithHole”** 45

Geometry - Collections

- MultiPoint, MultiLineString, MultiPolygon, GeometryCollection
- Useful spatial function
 - `ST_NumGeometries(geometry)` returns the number of parts in the collection
 - `ST_GeometryN(geometry,n)` returns the specified part

```
postgis_lab=# SELECT name, ST_AsText(geom), ST_NumGeometries(geom)
FROM geometries
```

```
WHERE name = 'Collection';
```

name	st_astext	st_numgeometries
Collection	GEOMETRYCOLLECTION(POINT(2 0),POLYGON((0 0,1 0,1 1,0 1,0 0)))	2

(1 row)

Geometry - Exercise

- What is the name and length of street which contains most points ?


(Table: nyc_streets)

- What is the area of each borough?

(Table: nyc_census_blocks)

Spatial Relationships


- **ST_Intersects(geometry A, geometry B):** returns TRUE if the two shapes have any space in common




```
postgis_lab=# SELECT name, boroname
postgis_lab=# FROM nyc_neighborhoods
postgis_lab=# WHERE ST_Intersects( geom,
postgis_lab=# (SELECT geom FROM nyc_subway_stations WHERE name = 'Broad St'));
      name      | boroname
-----+-----
Financial District | Manhattan
(1 row)
```

Spatial Relationships

- **ST_Touches(geometry A, geometry B):** tests whether two geometries touch at their boundaries, but do not intersect in their interiors




Point & Linestring



Multipoint & Linestring




Linestring & Linestring



Linestring & Polygon



Point & Polygon



Multipoint & Polygon

Spatial Relationships


- **ST_Within(geometry A, geometry B):** returns TRUE if the first geometry is completely within the second geometry
- **ST_Contains(geometry A, geometry B):** returns TRUE if the second geometry is completely contained by the first geometry.



```
postgis=# SELECT name
postgis=# FROM nyc_neighborhoods
postgis=# WHERE ST_Contains(geom, ST_GeomFromText('POINT(583571 4506714)',26918));
name
-----
Financial District
```

Spatial Relationships

- **ST_Distance(geometry A, geometry B):** calculates the shortest distance between two geometries.




```
postgres=# SELECT ST_Distance(
postgres=# ST_GeometryFromText('POINT(2 5)'),
postgres=# ST_GeometryFromText('LINESTRING(-2 2, 2 2)'));

 st_distance
-----
3
(1 row)
```

Spatial Relationships

- **ST_DWithin(geometry A, geometry B, radius):** Returns true if the geometries are within the specified distance (radius) of one another.



```
postgis_lab=# SELECT name FROM nyc_streets
postgis_lab=# WHERE ST_DWithin(geom,
postgis_lab=# (SELECT geom FROM nyc_streets WHERE name = 'Atlantic Commons'),
postgis_lab=# 10);
```

```
      name
-----
S Oxford St
Atlantic Commons
Cumberland St
(3 rows)
```

Spatial Relationships - Exercise

- Street 'Atlantic Commons'
- What neighborhood is Atlantic Commons in?
(Table:nyc_neighborhoods)
- Approximately how many people live on (within 50 meters of) Atlantic Commons?
(Table: nyc_census_blocks)

Spatial Relationships – Problem

- Find the neighborhood which contains Broad Station.
- **SELECT** name, boroname
FROM nyc_neighborhoods
WHERE ST_Contains(geom,
 (SELECT geom **FROM** nyc_subway_stations
 WHERE name = 'Broad St'));
- How about that if we want to select data from both tables?

Spatial Join

- Spatial Join allow you to combine information from different tables by using spatial relationships as the join key.
- Retrieve the station, neighborhood which contains Broad Station.

```
SELECT subways.name AS subway_name,  
        neighborhoods.name AS neighborhood_name,  
        neighborhoods.borname AS borough  
FROM nyc_neighborhoods AS neighborhoods  
JOIN nyc_subway_stations AS subways  
ON ST_Contains(neighborhoods.geom, subways.geom)  
WHERE subways.name = 'Broad St';
```

Spatial Join

- What is the population and racial make-up of each neighborhoods of Manhattan?
- **SELECT**
neighborhoods.name **AS** neighborhood_name,
Sum(cen.popn_total) **AS** population,
100.0 * Sum(cen.popn_white) / Sum(cen.popn_total) **AS** white_pct,
100.0 * Sum(cen.popn_black) / Sum(cen.popn_total) **AS** black_pct
FROM nyc_neighborhoods **AS** neighborhoods
JOIN nyc_census_blocks **AS** cen
ON ST_Intersects(neighborhoods.geom, cen.geom)
WHERE neighborhoods.borname = 'Manhattan'
GROUP BY neighborhoods.name
ORDER BY white_pct **DESC**;

Spatial Join - Exercise

- Each subway station may have many routes passing. e.g. J-train, M-train, Z-train pass through Broad Station.
 - Find out the racial make-up of within 200 meters of the A-train line.
- (Table: nyc_subway_stations, nyc_census_blocks)

Reference

1. <http://www.postgresqltutorial.com/>